

Semantics-based Distributed I/O for mpiBLAST

P. Balaji^α, W. Feng^β, J. Archuleta^β, H. Lin^δ, R. Kettimuthu^α, R. Thakur^α and X. Ma^δ

^α Argonne National Laboratory

^β Virginia Tech

^δ North Carolina State University

Traditional Distributed I/O

Issues with Distributed I/O

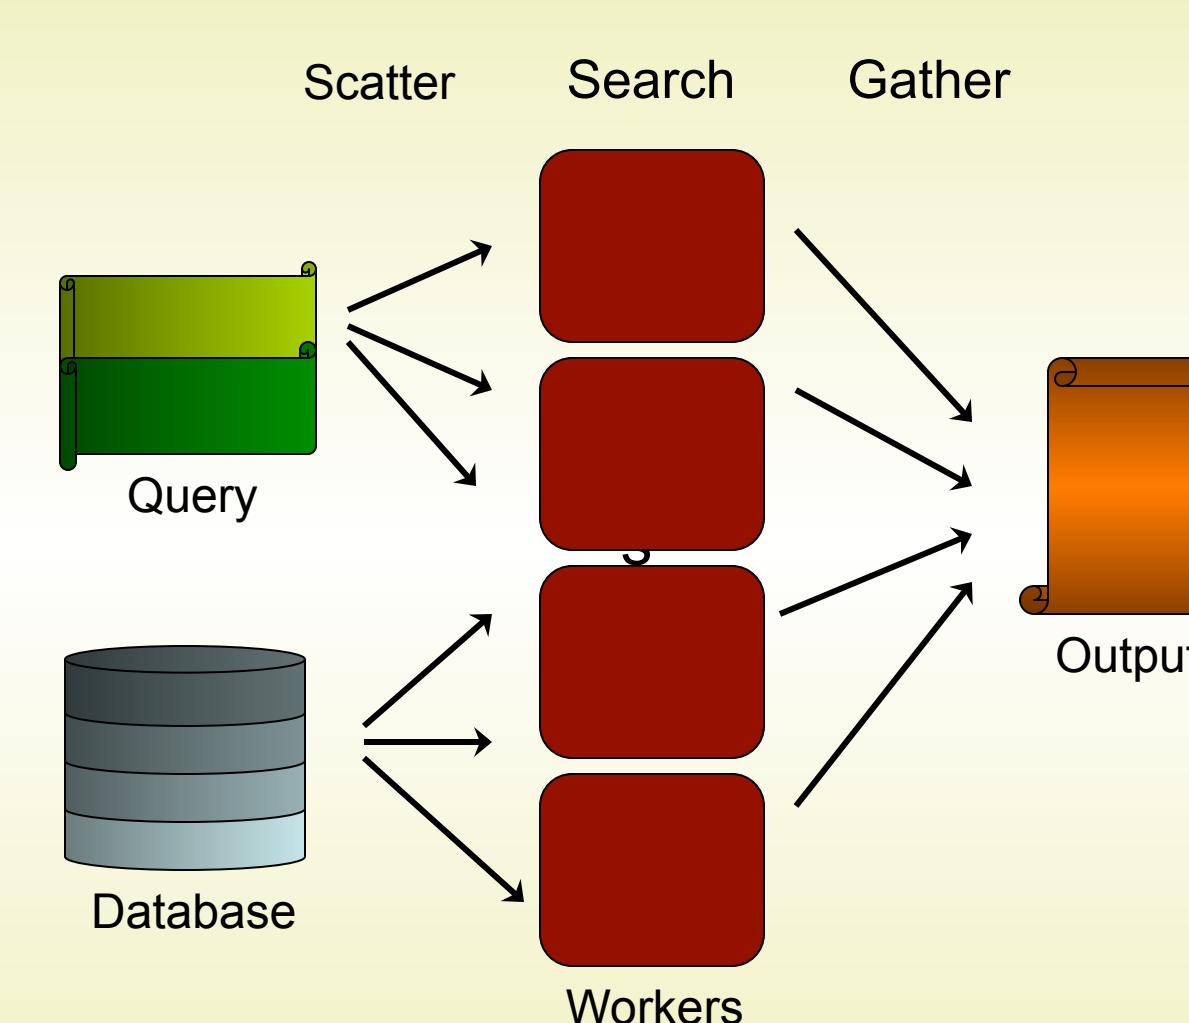
- High Latency
 - Synchronization operations heavily effected
- Low Bandwidth
 - High-bandwidth links are not accessible to everyone
- Data Encryption
 - Distributed I/O over the Internet might need to be encrypted in some environments
- And yet distributed I/O is essential
 - Large-scale computations might require resources that are not available at a single site
 - Scientists may need to access remote large-scale supercomputers which are not available locally

NSF TeraGrid

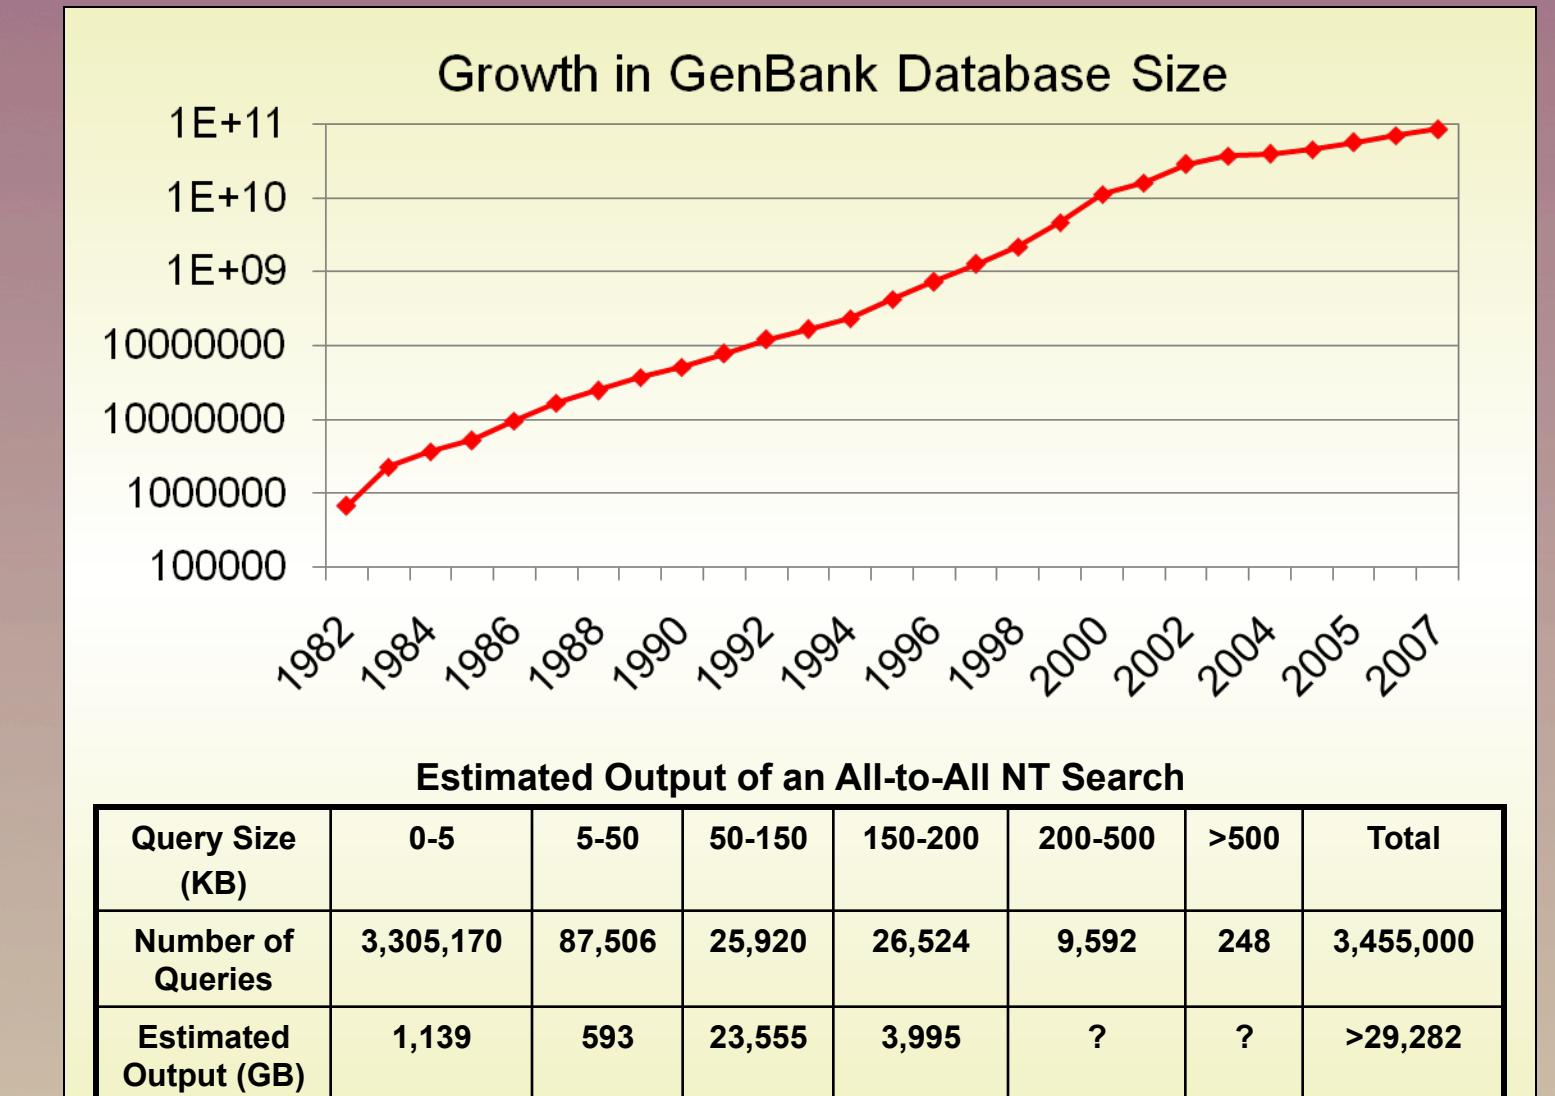


Sequence Search with mpiBLAST

mpiBLAST Working Model

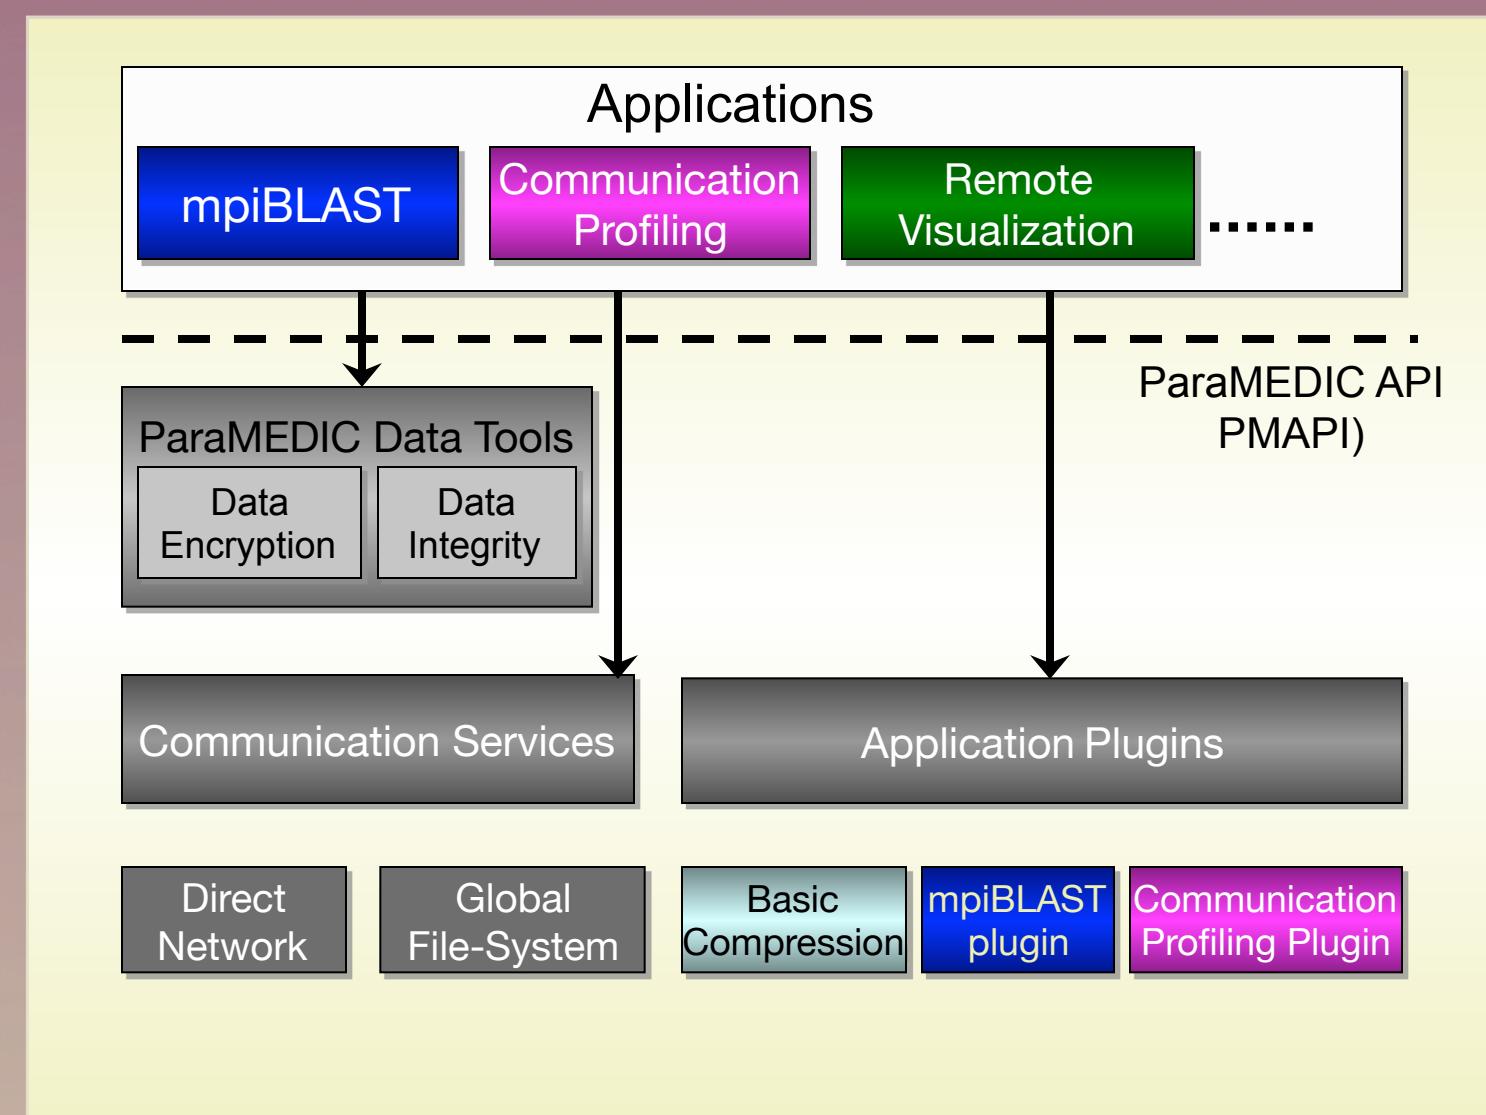


Distributed I/O in mpiBLAST



ParaMEDIC Framework

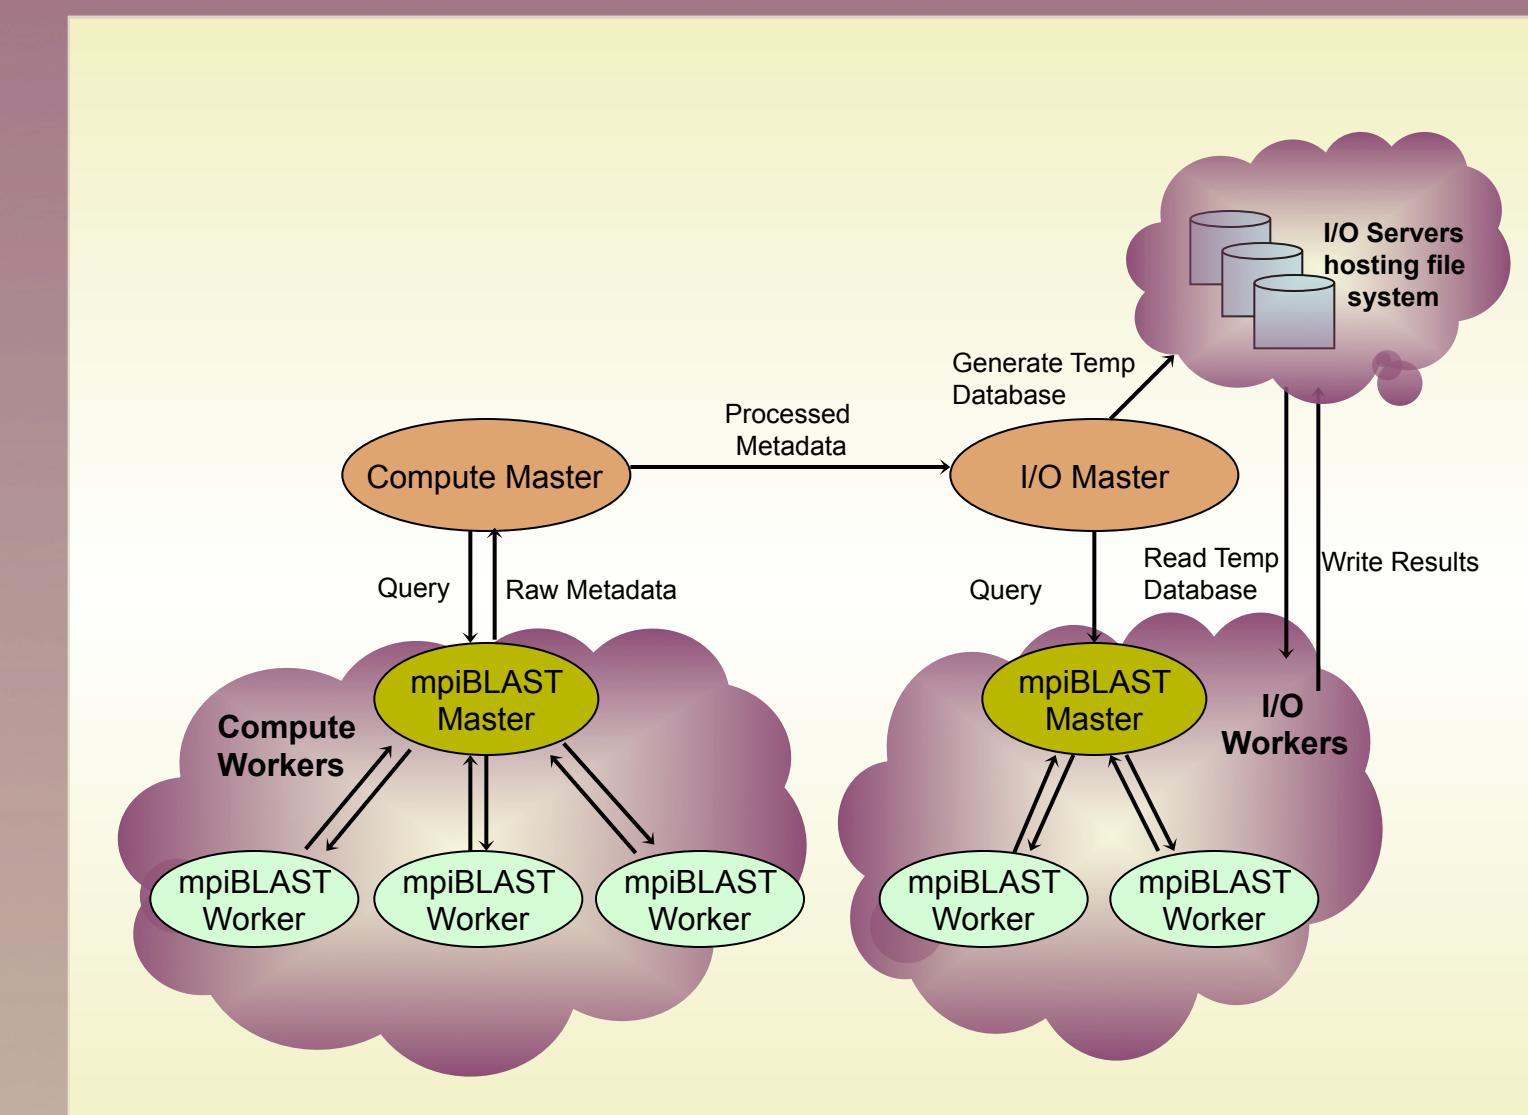
ParaMEDIC Architecture



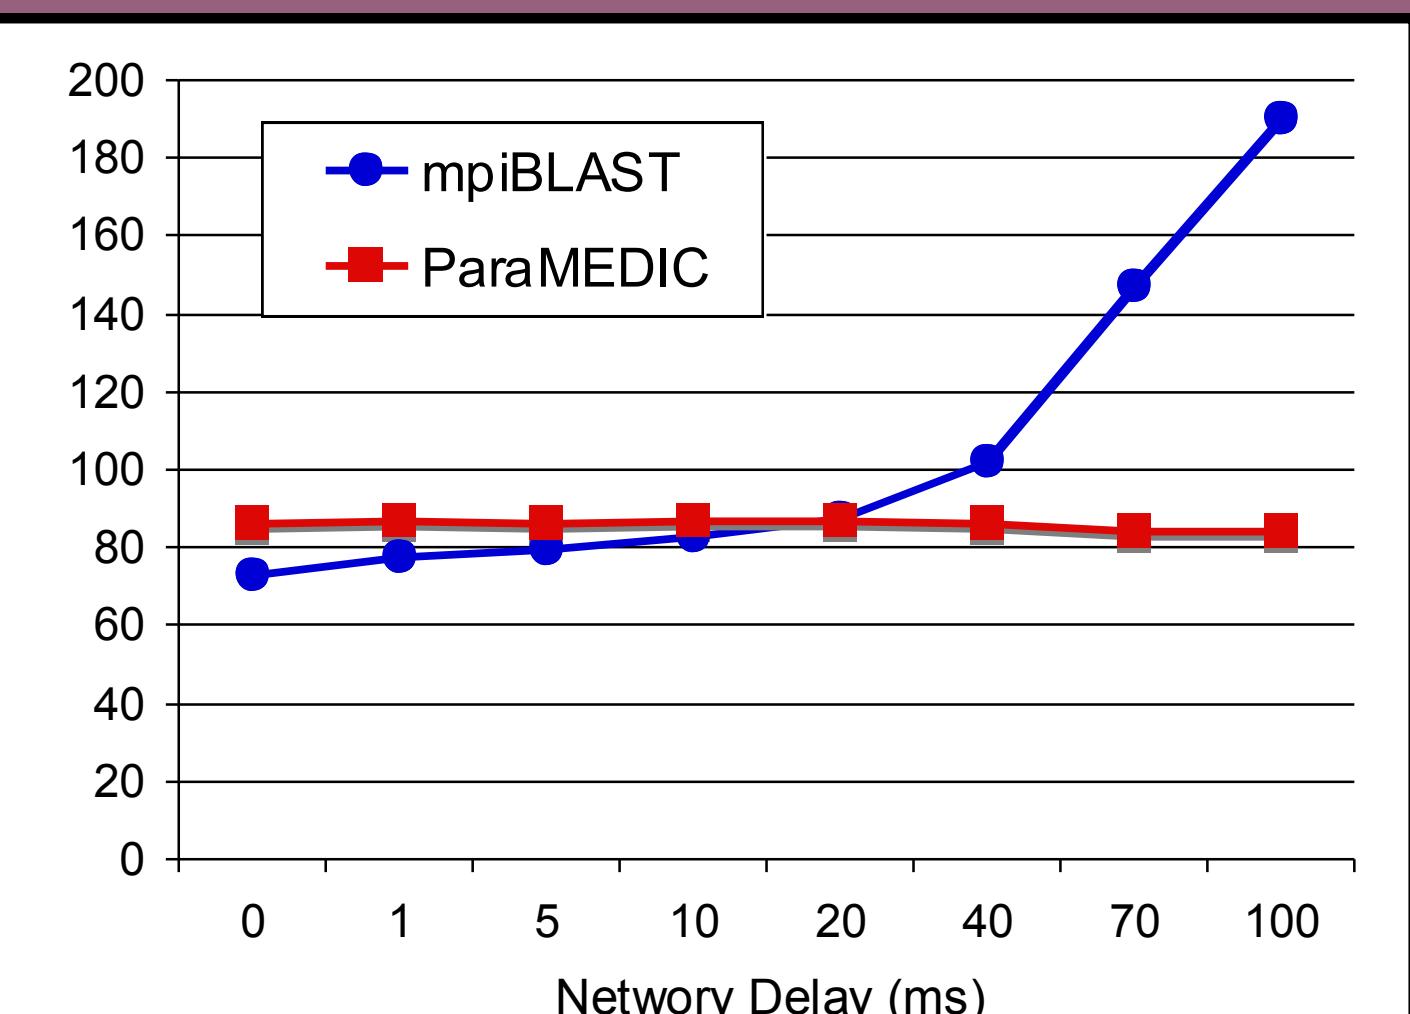
Understanding ParaMEDIC

- Primary Idea:
 - Data transformation, but not at the byte-level (unlike compression techniques)
 - Application specifies the appropriate description of the data → allows ParaMEDIC to process output as high-level objects, and not a stream of bytes
- With respect to mpiBLAST
 - Overall output is just a concatenation of matches and other descriptive information for each query sequence
 - Match scores, alignment information, etc.
 - Finding the database matches for each sequence is the compute intensive portion → needs to be stored
 - Rest can be discarded and regenerated

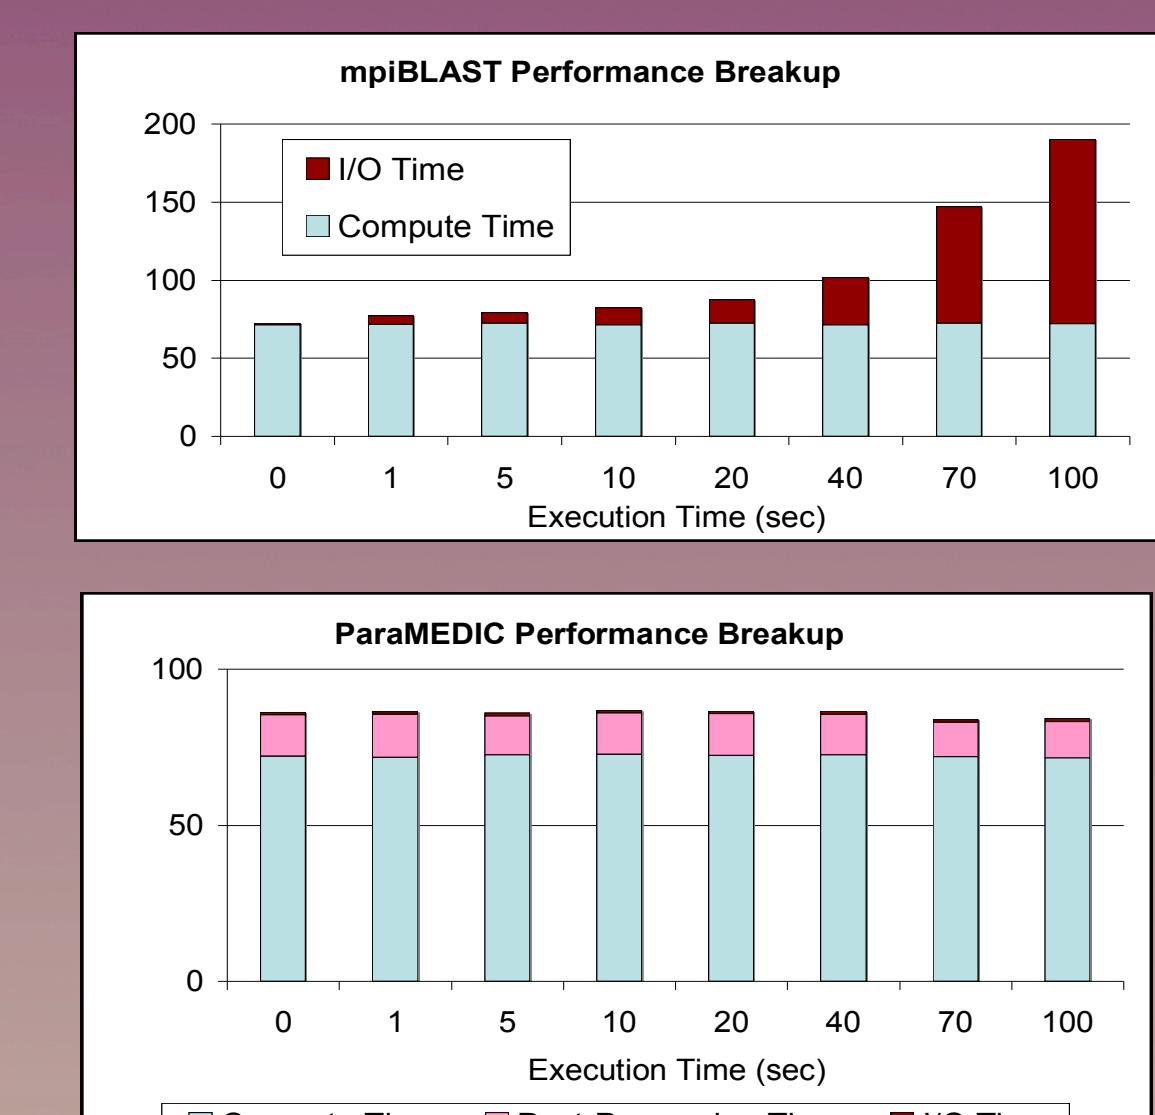
ParaMEDIC-powered mpiBLAST



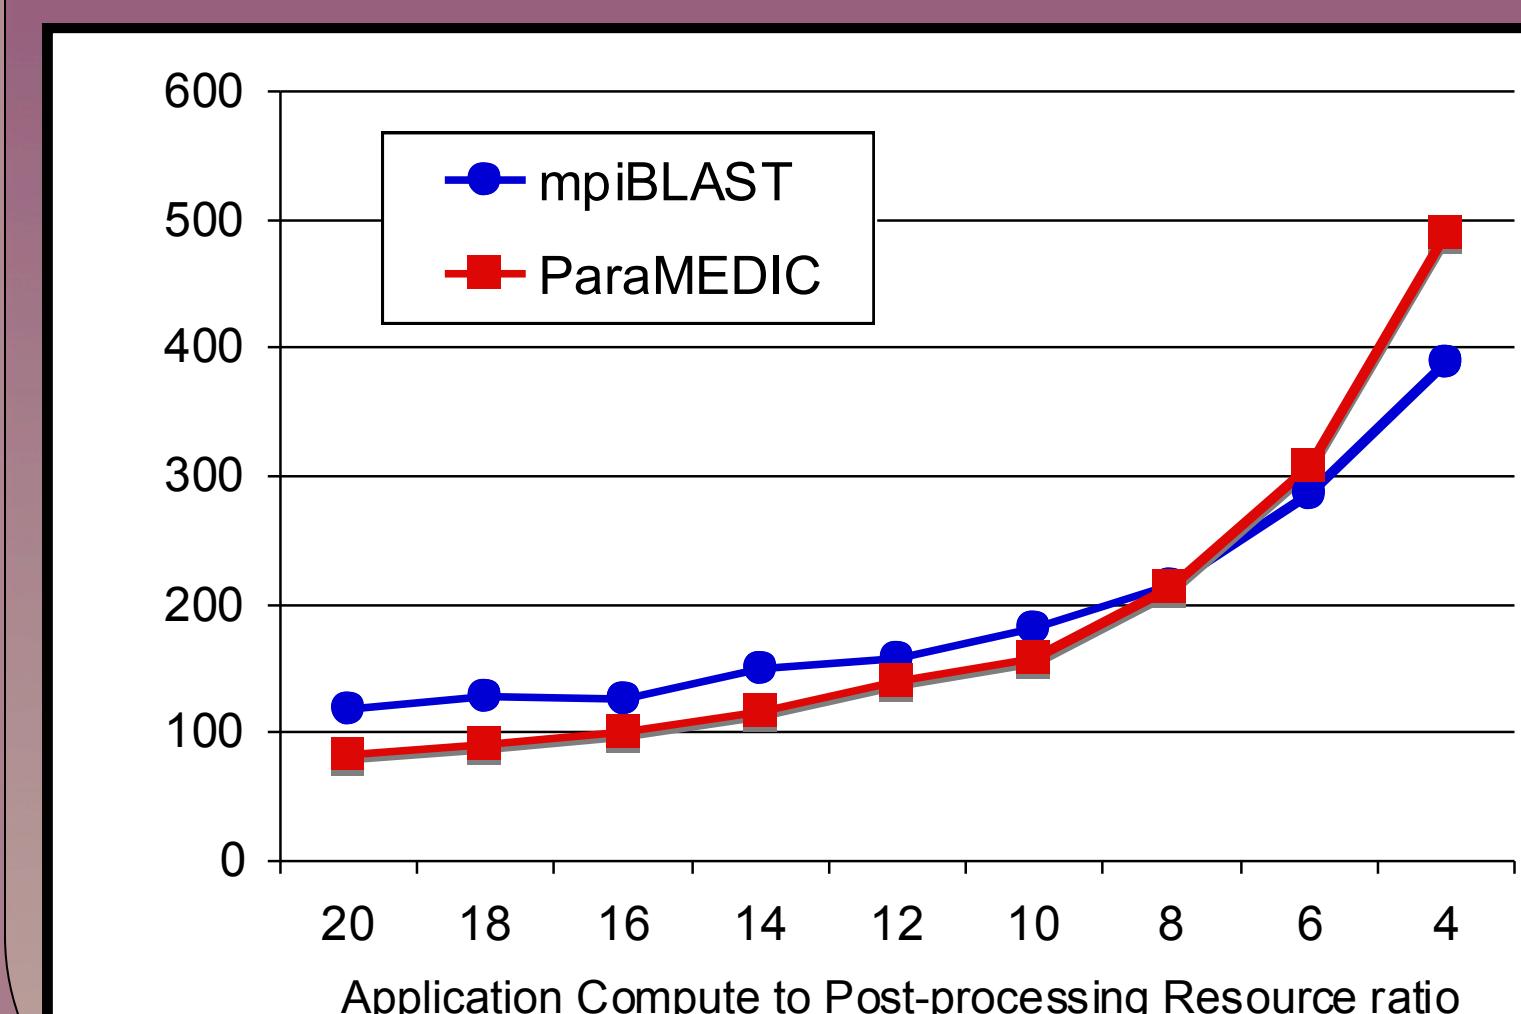
Impact of Network Latency



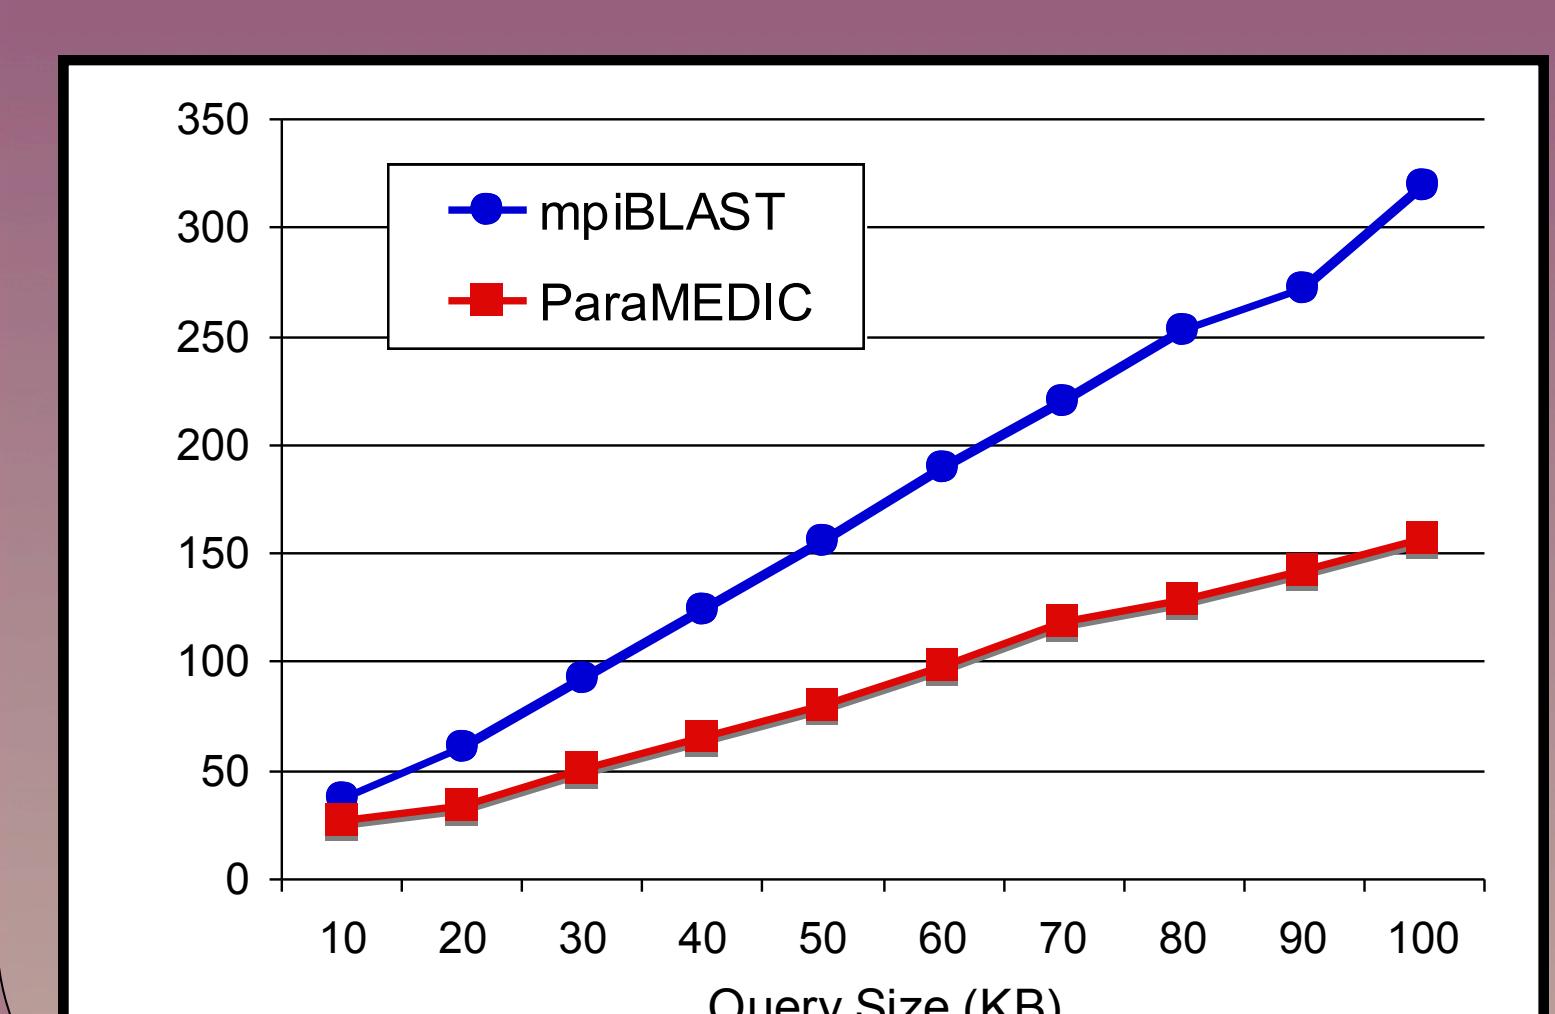
Network Delay: Performance Breakup



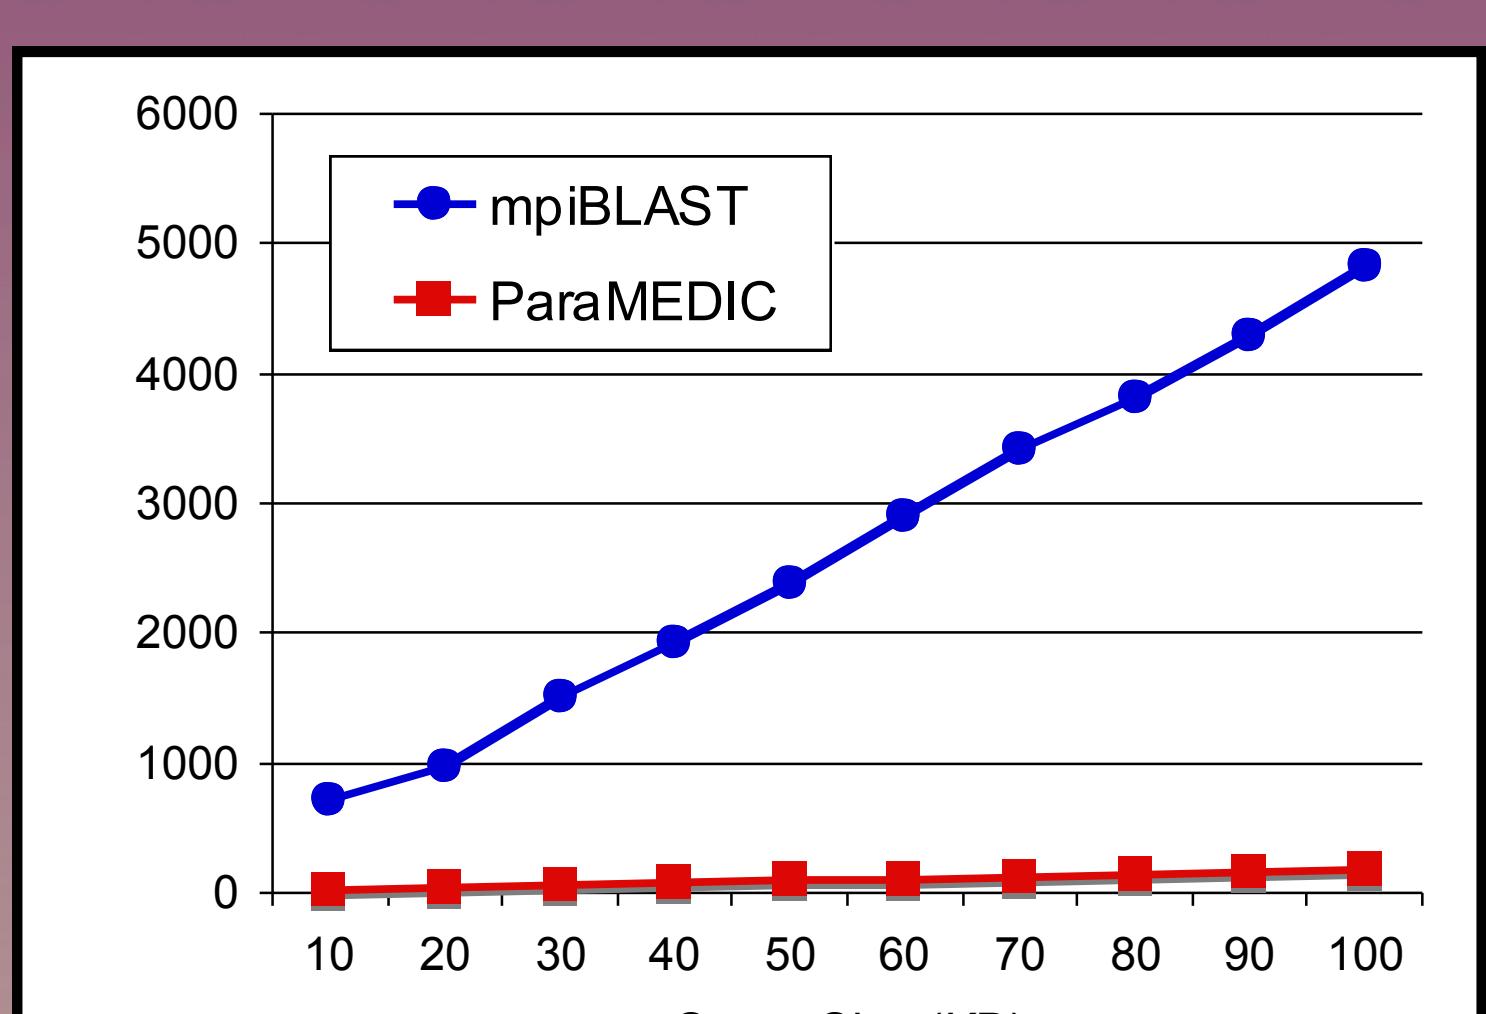
Trading Computation and IO



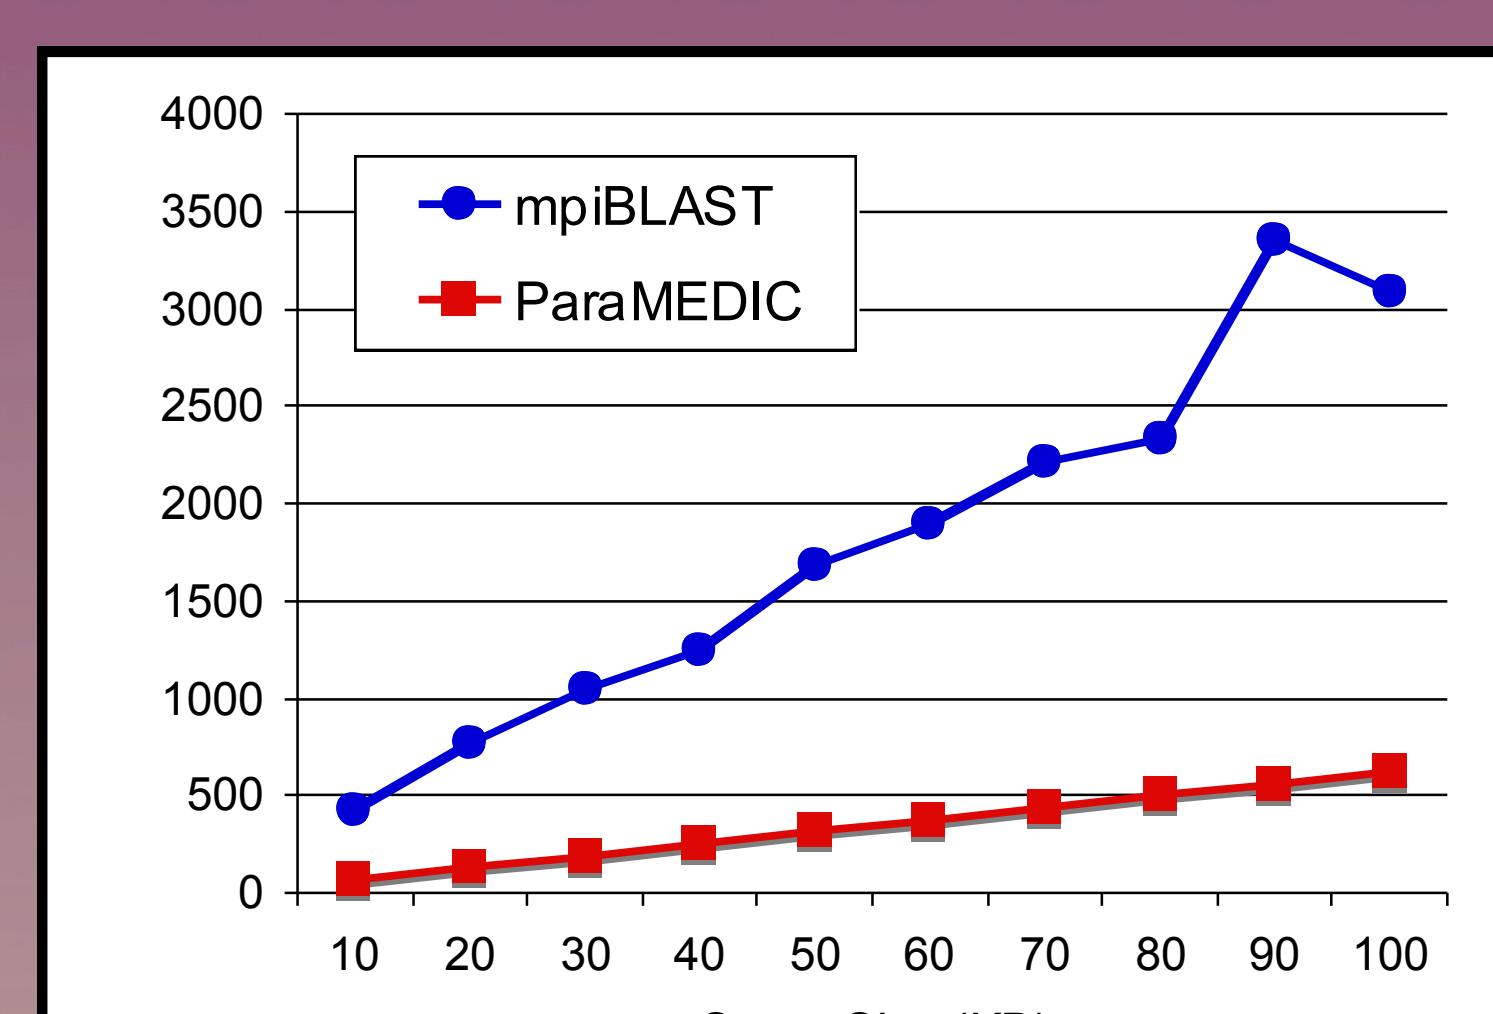
Impact of Encrypted File-Systems



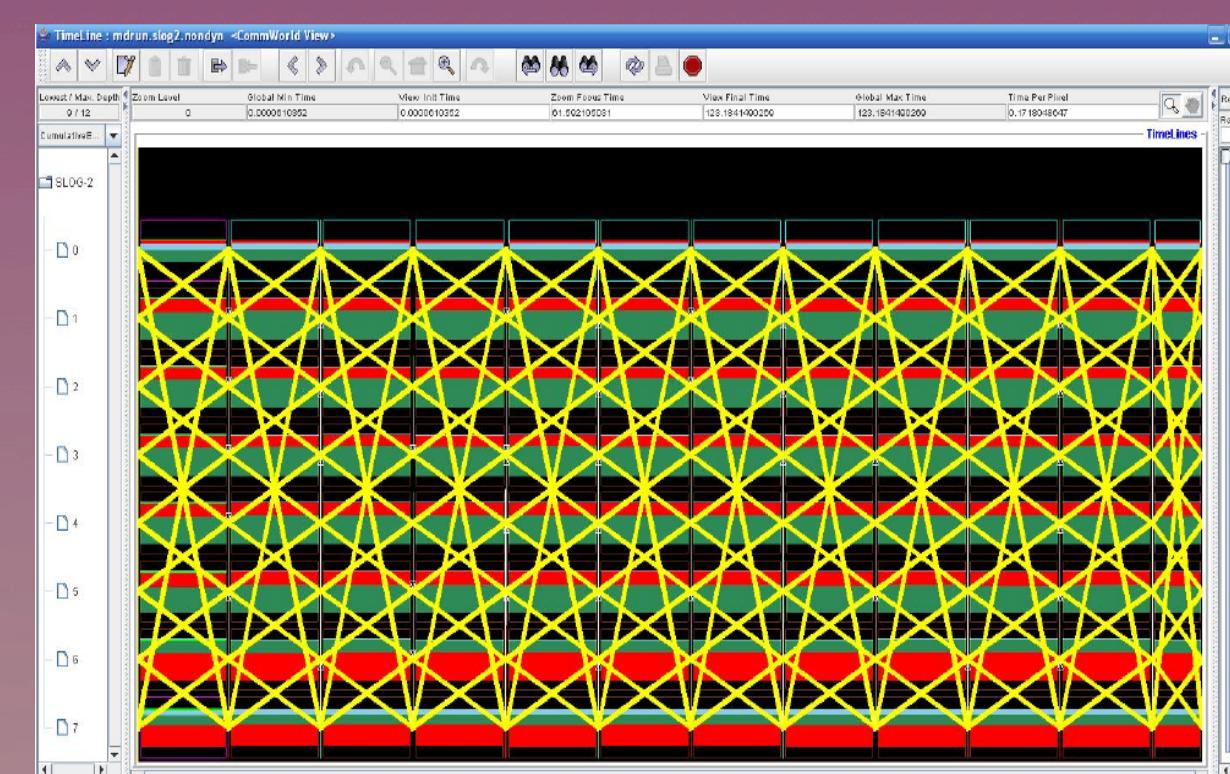
Argonne-VT Distributed Setup



Performance on TeraGrid



Other Applications: MPE Communication Profiler



- Scientific applications have periodic communication profiles
- ParaMEDIC uses FFT to find periodicity and process output
 - Preliminary results show 2-5X reduction in output

Conclusion

- Distributed I/O is a necessary evil
 - Large applications need resources from multiple centers
 - Scientists might need to use remote large-scale resources
- Impacted by several issues
 - High Latency, Low Bandwidth, Encryption requirements
- We propose the ParaMEDIC framework
 - Semantics-based Distributed I/O mechanism
 - ParaMEDIC uses application-specific plugins to "understand" what the output data means and process it
 - Output data converted to application-specific metadata, transported, and then converted back
- Trades a small amount of additional computation for potentially large benefits in I/O